

Classification and Diagnostic Tests

Last modified: 2026-10-09 01:12:45 (PDT)

Introduction

Classification is a core problem in statistics and machine learning: we seek to assign individuals or observations to one of several discrete categories based on available data. In medicine and epidemiology, classification problems arise constantly; for example, a clinician decides whether a patient has a disease based on evidence such as:

- test results;
- biomarkers;
- clinical signs.

A test can look highly accurate in isolation, yet its predictive value for an individual patient depends heavily on the prevalence of the condition in the population being tested. Understanding this interplay requires [Bayes' theorem](#) and the [law of total probability](#).

Definition 0.1 (Classification). A **classification problem** is a statistical problem in which we seek to assign each observation to one of two or more discrete categories (**classes**) based on its observed features or predictors x , using a **classification rule** c that maps features to classes:

$$c : x \mapsto c(x) \in \{1, \dots, K\} \tag{1}$$

In the binary case ($K = 2$), the two classes are often labeled “positive” and “negative”, or “diseased” and “healthy”.

Example 0.1 (A diagnostic test as a classification rule). A COVID-19 test assigns each person tested to one of $K = 2$ classes, “has COVID-19” or “does not have COVID-19”, based on a single feature x : the test’s result, positive (+) or negative ($\neg +$). The test is the classification rule (Definition 0.1) that assigns a positive result to “has COVID-19” and a negative result to “does not have COVID-19”:

$$c(x) \stackrel{\text{def}}{=} \begin{cases} \text{has COVID-19} & \text{if } x = + \\ \text{does not have COVID-19} & \text{if } x = \neg+ \end{cases} \quad (2)$$

Diagnostic test characteristics

Definition 0.2 (Sensitivity). The **sensitivity** of a diagnostic test is the probability that the test is positive (+), given that the person tested has the disease (D):

$$\text{sensitivity} \stackrel{\text{def}}{=} \Pr(+ \mid D) \quad (3)$$

Example 0.2 (Sensitivity of a COVID-19 test). Suppose a COVID-19 test is positive for 99% of the people who have COVID-19. Let D be the event “the person has COVID-19” and $+$ the event “the test is positive”. Then the test’s sensitivity (Definition 0.2) is:

$$\Pr(+ \mid D) = 0.99 \quad (4)$$

Definition 0.3 (Specificity). The **specificity** of a diagnostic test is the probability that the test is negative ($\neg+$), given that the person tested does not have the disease ($\neg D$):

$$\text{specificity} \stackrel{\text{def}}{=} \Pr(\neg+ \mid \neg D) \quad (5)$$

Example 0.3 (Specificity of a COVID-19 test). Suppose the COVID-19 test of Example 0.2 is negative for 99% of the people who do not have COVID-19. Then its specificity (Definition 0.3) is:

$$\Pr(\neg+ \mid \neg D) = 0.99 \quad (6)$$

By the [complement rule](#), its false positive rate is:

$$\begin{aligned} \Pr(+ \mid \neg D) &= 1 - \Pr(\neg+ \mid \neg D) && \text{(complement rule, conditional on } \neg D) \\ &= 1 - 0.99 && \text{(substitute the specificity)} \\ &= 0.01 && \text{(subtract)} \end{aligned} \quad (7)$$

Definition 0.4 (Prevalence). The **prevalence** of a disease in a population is the probability that a person drawn from that population has the disease (D):

$$\text{prevalence} \stackrel{\text{def}}{=} \Pr(D) \quad (8)$$

Example 0.4 (Prevalence of COVID-19). Suppose 7% of the population being tested with the COVID-19 test of Example 0.2 has COVID-19. Then the prevalence (Definition 0.4) is:

$$\Pr(D) = 0.07 \quad (9)$$

and, by the [complement rule](#), the probability that a person drawn from that population does not have COVID-19 is:

$$\begin{aligned} \Pr(\neg D) &= 1 - \Pr(D) && \text{(complement rule)} \\ &= 1 - 0.07 && \text{(substitute the prevalence)} \\ &= 0.93 && \text{(subtract)} \end{aligned} \quad (10)$$

Predictive values

Definition 0.5 (Positive predictive value (PPV)). The **positive predictive value** of a diagnostic test is the probability that the person tested has the disease (D), given that the test is positive (+):

$$\text{PPV} \stackrel{\text{def}}{=} \Pr(D \mid +) \quad (11)$$

Example 0.5 (PPV of a COVID-19 test). For the COVID-19 test of Example 0.2 and Example 0.3, used in the population of Example 0.4, [Bayes' theorem](#) gives the positive predictive value (Definition 0.5):

$$\Pr(D \mid +) \approx 0.88 \quad (12)$$

Even with a highly accurate test (99% sensitive and 99% specific), only about 88% of the people who test positive have COVID-19, because the prevalence (7%) is low enough that false positives make up a meaningful fraction of all positive tests.

Definition 0.6 (Negative predictive value (NPV)). The **negative predictive value** of a diagnostic test is the probability that the person tested does not have the disease ($\neg D$), given that the test is negative ($\neg +$):

$$\text{NPV} \stackrel{\text{def}}{=} \Pr(\neg D \mid \neg +) \quad (13)$$

Example 0.6 (NPV of a COVID-19 test). For the COVID-19 test of Example 0.2 and Example 0.3, used in the population of Example 0.4, the probability of a negative test for a person with COVID-19 is, by the [complement rule](#):

$$\begin{aligned}\Pr(\neg+ \mid D) &= 1 - \Pr(+ \mid D) && \text{(complement rule, conditional on } D) \\ &= 1 - 0.99 && \text{(substitute the sensitivity)} \\ &= 0.01 && \text{(subtract)}\end{aligned}\tag{14}$$

By [Bayes' theorem](#), with the denominator expanded by the [law of total probability](#) over the partition $\{D, \neg D\}$, the negative predictive value (Definition 0.6) is:

$$\begin{aligned}\Pr(\neg D \mid \neg+) &= \frac{\Pr(\neg+ \mid \neg D) \cdot \Pr(\neg D)}{\Pr(\neg+)} && \text{(Bayes' theorem)} \\ &= \frac{\Pr(\neg+ \mid \neg D) \cdot \Pr(\neg D)}{\Pr(\neg+ \mid \neg D) \cdot \Pr(\neg D) + \Pr(\neg+ \mid D) \cdot \Pr(D)} && \text{(law of total probability)} \\ &= \frac{0.99 \cdot 0.93}{0.99 \cdot 0.93 + 0.01 \cdot 0.07} && \text{(substitute the given values)} \\ &= \frac{0.9207}{0.9207 + 0.0007} && \text{(multiply each term in the numerator and denominator by } 0.0007) \\ &= \frac{0.9207}{0.9214} && \text{(add the denominator's two terms)} \\ &\approx 0.9992 && \text{(divide)}\end{aligned}\tag{15}$$

A negative result from this test is very reliable: only about 8 in 10,000 people who test negative have COVID-19.

PPV in terms of sensitivity, specificity, and prevalence

Exercise 0.1 (Probability of a positive test). Write the probability $\Pr(+)$ that a diagnostic test is positive in terms of the test's [sensitivity](#), [specificity](#), and the [prevalence](#) of the disease.

Solution

Solution 0.1. By the [law of total probability](#) over the partition $\{D, \neg D\}$:

$$\begin{aligned}
\Pr(+) &= \Pr(+ \mid D) \cdot \Pr(D) + \Pr(+ \mid \neg D) \cdot \Pr(\neg D) && \text{(law of total probability)} \\
&= \text{sensitivity} \cdot \Pr(D) + \Pr(+ \mid \neg D) \cdot \Pr(\neg D) && \text{(definition of sensitivity)} \\
&= \text{sensitivity} \cdot \text{prevalence} + \Pr(+ \mid \neg D) \cdot \Pr(\neg D) && \text{(definition of prevalence)} \\
&= \text{sensitivity} \cdot \text{prevalence} + (1 - \Pr(\neg + \mid \neg D)) \cdot \Pr(\neg D) && \text{(complement rule, conditional on } \neg D) \\
&= \text{sensitivity} \cdot \text{prevalence} + (1 - \text{specificity}) \cdot \Pr(\neg D) && \text{(definition of specificity)} \\
&= \text{sensitivity} \cdot \text{prevalence} + (1 - \text{specificity}) \cdot (1 - \Pr(D)) && \text{(complement rule)} \\
&= \text{sensitivity} \cdot \text{prevalence} + (1 - \text{specificity}) \cdot (1 - \text{prevalence}) && \text{(definition of prevalence)} \\
&&& (16)
\end{aligned}$$

Exercise 0.2 (PPV from Bayes' theorem). Write the [positive predictive value](#) of a diagnostic test in terms of its [sensitivity](#), [specificity](#), and the [prevalence](#) of the disease.

Solution

Solution 0.2. Abbreviate the sensitivity as *sens*, the specificity as *spec*, and the prevalence as *prev*. By [Bayes' theorem](#), with the probability of a positive test from Exercise 0.1:

$$\begin{aligned}
\text{PPV} &\stackrel{\text{def}}{=} \Pr(D \mid +) && \text{(definition of PPV)} \\
&= \frac{\Pr(+ \mid D) \cdot \Pr(D)}{\Pr(+)} && \text{(Bayes' theorem)} \\
&= \frac{\text{sens} \cdot \Pr(D)}{\Pr(+)} && \text{(definition of sensitivity)} \\
&= \frac{\text{sens} \cdot \text{prev}}{\Pr(+)} && \text{(definition of prevalence)} \\
&= \frac{\text{sens} \cdot \text{prev}}{\text{sens} \cdot \text{prev} + (1 - \text{spec}) \cdot (1 - \text{prev})} && \text{(probability of a positive test)} \\
&&& (17)
\end{aligned}$$

Exercise 0.3 (PPV in terms of two ratios). Rewrite the expression for the positive predictive value from Exercise 0.2 so that the sensitivity, specificity, and prevalence appear only through the ratio $\frac{1-\text{spec}}{\text{sens}}$ of the false positive rate to the sensitivity and the ratio $\frac{1-\text{prev}}{\text{prev}}$ of non-diseased to diseased people in the population.

Solution

Solution 0.3. Write $a \stackrel{\text{def}}{=} \text{sens} \cdot \text{prev}$ and $b \stackrel{\text{def}}{=} (1 - \text{spec}) \cdot (1 - \text{prev})$, so that Exercise 0.2 reads $\text{PPV} = \frac{a}{a+b}$, with $a > 0$ whenever the sensitivity and prevalence are positive. Then:

$$\begin{aligned}
PPV &= \frac{a}{a+b} && \text{(PPV from Bayes' theorem)} \\
&= \frac{1}{\left(\frac{a+b}{a}\right)} && \text{(a fraction equals one over its reciprocal, since } a > 0\text{)} \\
&= \frac{1}{\frac{a}{a} + \frac{b}{a}} && \text{(split the fraction over the sum in its numerator)} \\
&= \frac{1}{1 + \frac{b}{a}} && \left(\frac{a}{a} = 1\right) \\
&= \frac{1}{1 + \frac{(1-\text{spec}) \cdot (1-\text{prev})}{\text{sens} \cdot \text{prev}}} && \text{(substitute the definitions of } a \text{ and } b\text{)} \\
&= \frac{1}{1 + \frac{1-\text{spec}}{\text{sens}} \cdot \frac{1-\text{prev}}{\text{prev}}} && \text{(a quotient of products is the product of the quotients)}
\end{aligned} \tag{18}$$

Theorem 0.1 (PPV in terms of sensitivity, specificity, and prevalence). *The positive predictive value of a diagnostic test with positive sensitivity, used in a population with positive prevalence, depends on its sensitivity, its specificity, and the prevalence of the disease only through the ratio of the false positive rate to the sensitivity and the ratio of non-diseased to diseased people in the population:*

$$PPV = \frac{1}{1 + \frac{1-\text{spec}}{\text{sens}} \cdot \frac{1-\text{prev}}{\text{prev}}} \tag{19}$$

i Proof

Proof. This result is the solution to Exercise 0.3, which builds on Exercise 0.2 and Exercise 0.1. □

Example 0.7 (PPV of a COVID-19 test, from the two ratios). For the COVID-19 test of Example 0.2 and Example 0.3, used in the population of Example 0.4, Theorem 0.1 gives:

$$\begin{aligned}
\text{PPV} &= \frac{1}{1 + \frac{1-0.99}{0.99} \cdot \frac{1-0.07}{0.07}} && \text{(PPV in terms of the two ratios, with the given values)} \\
&= \frac{1}{1 + \frac{0.01}{0.99} \cdot \frac{0.93}{0.07}} && \text{(subtract in each numerator)} \\
&\approx \frac{1}{1 + 0.0101 \cdot 13.29} && \text{(divide in each ratio)} \\
&\approx \frac{1}{1 + 0.134} && \text{(multiply)} \\
&= \frac{1}{1.134} && \text{(add)} \\
&\approx 0.88 && \text{(divide)}
\end{aligned} \tag{20}$$

This value matches Example 0.5.

If the same test were used in a population with prevalence 0.1% instead, the ratio of non-diseased to diseased people would be $0.999/0.001 = 999$, and the PPV would fall to

$$\begin{aligned}
\text{PPV} &= \frac{1}{1 + \frac{0.01}{0.99} \cdot 999} && \text{(PPV in terms of the two ratios, with the new prevalence)} \\
&\approx \frac{1}{1 + 0.0101 \cdot 999} && \text{(divide)} \\
&\approx \frac{1}{1 + 10.09} && \text{(multiply)} \\
&= \frac{1}{11.09} && \text{(add)} \\
&\approx 0.09 && \text{(divide)}
\end{aligned} \tag{21}$$

so about 9 in 10 positive results would be false positives.

Remark 0.1 (The PPV depends on the population, not only on the test). The sensitivity and specificity describe the test itself, but the ratio $\frac{1-\text{prev}}{\text{prev}}$ in Theorem 0.1 describes the population being tested. When the disease is rare, that ratio is large, so even a small false positive rate can make the PPV low, as Example 0.7 shows.

Agreement between two classifiers

So far, we have compared a test with the true disease status. The same tools apply whenever we compare two classifiers that label the same observations: for example, a new diagnostic test against a reference test, or an automated grader (an “agent”) against a human grader. We call

the classifier being evaluated the **agent** and the one we compare it with the **reference**.

Definition 0.7 (Agreement table). For n observations that an agent and a reference each classify as positive or negative, the **agreement table** counts the four combinations:

	Reference positive	Reference negative	Total
Agent positive	a	b	$a + b$
Agent negative	c	d	$c + d$
Total	$a + c$	$b + d$	n

The cells a and d count agreements, the cells b and c count disagreements, and the four cells together count every observation:

$$n \stackrel{\text{def}}{=} a + b + c + d \quad (22)$$

Example 0.8 (Agreement table for a COVID-19 test). Apply the COVID-19 test of Example 0.2 and Example 0.3 to $n = 10,000$ people from the population of Example 0.4, with the true disease status as the reference. With 7% prevalence, $0.07 \times 10,000 = 700$ people have COVID-19 and $10,000 - 700 = 9,300$ do not. With 99% sensitivity, the test is positive for $0.99 \times 700 = 693$ of the 700; with 99% specificity, it is negative for $0.99 \times 9,300 = 9,207$ of the 9,300. So the test is positive for $9,300 - 9,207 = 93$ people without COVID-19 and negative for $700 - 693 = 7$ people with it. The agreement table (Definition 0.7) is:

	COVID-19	No COVID-19	Total
Test positive	$a = 693$	$b = 93$	786
Test negative	$c = 7$	$d = 9,207$	9,214
Total	700	9,300	10,000

```
agree_covid <- c(a = 693, b = 93, c = 7, d = 9207)
agree_covid
#>    a    b    c    d
#> 693   93    7 9207
```

Agreement proportions

Definition 0.8 (Overall agreement). The **overall agreement** of an agreement table (Definition 0.7) is the proportion of observations on which the agent and the reference agree:

$$p_o \stackrel{\text{def}}{=} \frac{a + d}{n} \quad (23)$$

Example 0.9 (Overall agreement for the COVID-19 test). For the agreement table of Example 0.8, the overall agreement (Definition 0.8) is:

$$\begin{aligned} p_o &= \frac{693 + 9,207}{10,000} && \text{(definition of } p_o, \text{ with the table's counts)} \\ &= \frac{9,900}{10,000} && \text{(add)} \\ &= 0.99 && \text{(divide)} \end{aligned} \quad (24)$$

Definition 0.9 (Positive agreement). The **positive agreement** (also called the *average positive agreement*) of an agreement table (Definition 0.7) with $2a + b + c > 0$ is the number of positive calls the agent and the reference share, counted once for each classifier, as a proportion of all their positive calls:

$$p_{\text{pos}} \stackrel{\text{def}}{=} \frac{2a}{2a + b + c} \quad (25)$$

Source

- 1.

Example 0.10 (Positive agreement for the COVID-19 test). For the agreement table of Example 0.8, the positive agreement (Definition 0.9) is:

$$\begin{aligned} p_{\text{pos}} &= \frac{2 \times 693}{2 \times 693 + 93 + 7} && \text{(definition of } p_{\text{pos}}, \text{ with the table's counts)} \\ &= \frac{1,386}{1,386 + 93 + 7} && \text{(multiply)} \\ &= \frac{1,386}{1,486} && \text{(add)} \\ &\approx 0.933 && \text{(divide)} \end{aligned} \quad (26)$$

The positive agreement is lower than the overall agreement (Example 0.9), because the 93

false positives are large relative to the 693 true positives.

Definition 0.10 (Negative agreement). The **negative agreement** (also called the *average negative agreement*) of an agreement table (Definition 0.7) with $2d + b + c > 0$ is the number of negative calls the agent and the reference share, counted once for each classifier, as a proportion of all their negative calls:

$$p_{\text{neg}} \stackrel{\text{def}}{=} \frac{2d}{2d + b + c} \quad (27)$$

Source

1.

Example 0.11 (Negative agreement for the COVID-19 test). For the agreement table of Example 0.8, the negative agreement (Definition 0.10) is:

$$\begin{aligned} p_{\text{neg}} &= \frac{2 \times 9,207}{2 \times 9,207 + 93 + 7} && \text{(definition of } p_{\text{neg}}, \text{ with the table's counts)} \\ &= \frac{18,414}{18,414 + 93 + 7} && \text{(multiply)} \\ &= \frac{18,414}{18,514} && \text{(add)} \\ &\approx 0.995 && \text{(divide)} \end{aligned} \quad (28)$$

The negative agreement is even higher than the overall agreement (Example 0.9), because the 100 disagreements are small relative to the 9,207 shared negative calls.

Definition 0.11 (Agent-only rate). The **agent-only rate** of an agreement table (Definition 0.7) is the proportion of observations that the agent calls positive and the reference calls negative:

$$p_{\text{agent}} \stackrel{\text{def}}{=} \frac{b}{n} \quad (29)$$

Example 0.12 (Agent-only rate for the COVID-19 test). For the agreement table of Example 0.8, the agent-only rate (Definition 0.11) is:

$$\begin{aligned} p_{\text{agent}} &= \frac{93}{10,000} && \text{(definition of } p_{\text{agent}}, \text{ with the table's counts)} \\ &= 0.0093 && \text{(divide)} \end{aligned} \quad (30)$$

so about 9 of every 1,000 people tested are called positive by the test and negative by the reference.

Definition 0.12 (Test measures estimated from an agreement table). Treating the reference as the truth, an agreement table (Definition 0.7) estimates four measures of the agent:

- its [sensitivity](#);
- its [specificity](#);
- its [positive predictive value](#);
- its [negative predictive value](#).

Each estimate is the matching proportion of the table's cells, defined when its denominator is positive:

$$\begin{aligned}\widehat{\text{sensitivity}} &\stackrel{\text{def}}{=} \frac{a}{a + c}, \\ \widehat{\text{specificity}} &\stackrel{\text{def}}{=} \frac{d}{b + d}, \\ \widehat{\text{PPV}} &\stackrel{\text{def}}{=} \frac{a}{a + b}, \\ \widehat{\text{NPV}} &\stackrel{\text{def}}{=} \frac{d}{c + d}.\end{aligned}\tag{31}$$

Example 0.13 (Test measures for the COVID-19 test). For the agreement table of Example 0.8, Definition 0.12 gives:

```
test_measures <- with(as.list(agree_covid), c(
  sensitivity = a / (a + c),
  specificity = d / (b + d),
  ppv = a / (a + b),
  npv = d / (c + d)
))
round(test_measures, 4)
#> sensitivity specificity      ppv      npv
#>      0.9900      0.9900  0.8817  0.9992
```

These estimates match the test's sensitivity (0.99) and specificity (0.99), its positive predictive value (about 0.88, Example 0.5), and its negative predictive value (about 0.999, Example 0.6), because the table's counts are the expected counts for those values.

Positive and negative agreement as harmonic means

Definition 0.13 (Harmonic mean). The **harmonic mean** of two positive numbers x and y is the reciprocal of the average of their reciprocals:

$$H(x, y) \stackrel{\text{def}}{=} \frac{2}{\frac{1}{x} + \frac{1}{y}} \quad (32)$$

Example 0.14 (Harmonic mean of 0.5 and 1). By Definition 0.13:

$$\begin{aligned} H(0.5, 1) &= \frac{2}{\frac{1}{0.5} + \frac{1}{1}} && \text{(definition of } H) \\ &= \frac{2}{2 + 1} && \text{(take each reciprocal)} \\ &= \frac{2}{3} && \text{(add)} \\ &\approx 0.667 && \text{(divide)} \end{aligned} \quad (33)$$

The ordinary average of 0.5 and 1 is 0.75; the harmonic mean sits closer to the smaller number.

Exercise 0.4 (Positive agreement from PPV and sensitivity). For an agreement table (Definition 0.7) with $a > 0$, write the **harmonic mean** of the estimated positive predictive value and sensitivity (Definition 0.12) in terms of the table's cells.

Solution

Solution 0.4. Because $a > 0$, both $a + b$ and $a + c$ are positive, so both estimates are defined and positive:

$$\begin{aligned}
H(\widehat{\text{PPV}}, \widehat{\text{sensitivity}}) &= \frac{2}{\frac{1}{\widehat{\text{PPV}}} + \frac{1}{\widehat{\text{sensitivity}}}} && \text{(definition of } H) \\
&= \frac{2}{\frac{1}{a/(a+b)} + \frac{1}{\widehat{\text{sensitivity}}}} && \text{(definition of } \widehat{\text{PPV}}) \\
&= \frac{2}{\frac{1}{a/(a+b)} + \frac{1}{a/(a+c)}} && \text{(definition of } \widehat{\text{sensitivity}}) \\
&= \frac{2}{\frac{a+b}{a} + \frac{1}{a/(a+c)}} && \text{(reciprocal of a fraction)} \\
&= \frac{2}{\frac{a+b}{a} + \frac{a+c}{a}} && \text{(reciprocal of a fraction)} \\
&= \frac{2}{\frac{(a+b)+(a+c)}{a}} && \text{(add fractions with a common denominator)} \\
&= \frac{2}{\frac{a+b+a+c}{a}} && \text{(remove parentheses)} \\
&= \frac{2}{\frac{a+a+b+c}{a}} && \text{(reorder terms)} \\
&= \frac{2}{\frac{2a+b+c}{a}} && (a + a = 2a) \\
&= 2 \cdot \frac{a}{2a+b+c} && \text{(dividing by a fraction multiplies by its reciprocal)} \\
&= \frac{2a}{2a+b+c} && \text{(multiply)}
\end{aligned} \tag{34}$$

Theorem 0.2 (Positive agreement is a harmonic mean). *For an agreement table with $a > 0$, the positive agreement (Definition 0.9) is the harmonic mean (Definition 0.13) of the estimated positive predictive value and sensitivity (Definition 0.12):*

$$p_{\text{pos}} = H(\widehat{\text{PPV}}, \widehat{\text{sensitivity}}) \tag{35}$$

i Proof

Proof. By Exercise 0.4, the harmonic mean equals $2a/(2a + b + c)$, which is p_{pos} by Definition 0.9. $\square$

Example 0.15 (Positive agreement of the COVID-19 test as a harmonic mean). For the agreement table of Example 0.8, the estimated positive predictive value is $693/786$ and the estimated sensitivity is $693/700$ (Example 0.13). Theorem 0.2 gives:

$$\begin{aligned}
 p_{\text{pos}} &= \frac{2}{\frac{1}{693/786} + \frac{1}{693/700}} && \text{(harmonic mean of the two estimates)} \\
 &= \frac{2}{\frac{786}{693} + \frac{1}{693/700}} && \text{(reciprocal of a fraction)} \\
 &= \frac{2}{\frac{786}{693} + \frac{700}{693}} && \text{(reciprocal of a fraction)} \\
 &= \frac{2}{\frac{1,486}{693}} && \text{(add fractions with a common denominator)} \\
 &= 2 \cdot \frac{693}{1,486} && \text{(dividing by a fraction multiplies by its reciprocal)} \\
 &= \frac{1,386}{1,486} && \text{(multiply)} \\
 &\approx 0.933 && \text{(divide)}
 \end{aligned} \tag{36}$$

which matches Example 0.10. The lower positive predictive value (about 0.88) pulls the positive agreement down from the sensitivity (0.99), as in Example 0.14.

Exercise 0.5 (Negative agreement from NPV and specificity). For an agreement table (Definition 0.7) with $d > 0$, write the **harmonic mean** of the estimated negative predictive value and specificity (Definition 0.12) in terms of the table's cells.

Solution

Solution 0.5. Because $d > 0$, both $c + d$ and $b + d$ are positive, so both estimates are defined and positive:

$$\begin{aligned}
H(\widehat{\text{NPV}}, \widehat{\text{specificity}}) &= \frac{2}{\frac{1}{\widehat{\text{NPV}}} + \frac{1}{\widehat{\text{specificity}}}} && \text{(definition of } H) \\
&= \frac{2}{\frac{1}{d/(c+d)} + \frac{1}{\widehat{\text{specificity}}}} && \text{(definition of } \widehat{\text{NPV}}) \\
&= \frac{2}{\frac{1}{d/(c+d)} + \frac{1}{d/(b+d)}} && \text{(definition of } \widehat{\text{specificity}}) \\
&= \frac{2}{\frac{c+d}{d} + \frac{1}{d/(b+d)}} && \text{(reciprocal of a fraction)} \\
&= \frac{2}{\frac{c+d}{d} + \frac{b+d}{d}} && \text{(reciprocal of a fraction)} \\
&= \frac{2}{\frac{(c+d)+(b+d)}{d}} && \text{(add fractions with a common denominator)} \\
&= \frac{2}{\frac{c+d+b+d}{d}} && \text{(remove parentheses)} \\
&= \frac{2}{\frac{d+d+b+c}{d}} && \text{(reorder terms)} \\
&= \frac{2}{\frac{2d+b+c}{d}} && (d + d = 2d) \\
&= 2 \cdot \frac{d}{2d + b + c} && \text{(dividing by a fraction multiplies by its reciprocal)} \\
&= \frac{2d}{2d + b + c} && \text{(multiply)}
\end{aligned} \tag{37}$$

Theorem 0.3 (Negative agreement is a harmonic mean). *For an agreement table with $d > 0$, the negative agreement (Definition 0.10) is the harmonic mean (Definition 0.13) of the estimated negative predictive value and specificity (Definition 0.12):*

$$p_{\text{neg}} = H(\widehat{\text{NPV}}, \widehat{\text{specificity}}) \tag{38}$$

i Proof

Proof. By Exercise 0.5, the harmonic mean equals $2d/(2d + b + c)$, which is p_{neg} by Definition 0.10. $\square$

Example 0.16 (Negative agreement of the COVID-19 test as a harmonic mean). For the agreement table of Example 0.8, the estimated negative predictive value is $9,207/9,214$ and the estimated specificity is $9,207/9,300$ (Example 0.13). Theorem 0.3 gives:

$$\begin{aligned}
 p_{\text{neg}} &= \frac{2}{\frac{1}{9,207/9,214} + \frac{1}{9,207/9,300}} && \text{(harmonic mean of the two estimates)} \\
 &= \frac{2}{\frac{9,214}{9,207} + \frac{1}{9,207/9,300}} && \text{(reciprocal of a fraction)} \\
 &= \frac{2}{\frac{9,214}{9,207} + \frac{9,300}{9,207}} && \text{(reciprocal of a fraction)} \\
 &= \frac{2}{\frac{18,514}{9,207}} && \text{(add fractions with a common denominator)} \\
 &= 2 \cdot \frac{9,207}{18,514} && \text{(dividing by a fraction multiplies by its reciprocal)} \\
 &= \frac{18,414}{18,514} && \text{(multiply)} \\
 &\approx 0.995 && \text{(divide)}
 \end{aligned} \tag{39}$$

which matches Example 0.11.

Cohen's kappa

Definition 0.14 (Chance agreement). The **chance agreement** of an agreement table (Definition 0.7) is the overall agreement we would expect if the agent and the reference classified independently, each with its observed proportion of positive calls:

$$p_e \stackrel{\text{def}}{=} \frac{(a+b)(a+c) + (c+d)(b+d)}{n^2} \tag{40}$$

Source

2.

Example 0.17 (Chance agreement for the COVID-19 test). For the agreement table of Example 0.8, the chance agreement (Definition 0.14) is:

$$\begin{aligned}
p_e &= \frac{786 \times 700 + 9,214 \times 9,300}{10,000^2} && \text{(definition of } p_e, \text{ with the table's counts)} \\
&= \frac{550,200 + 85,690,200}{10,000^2} && \text{(multiply)} \\
&= \frac{86,240,400}{10,000^2} && \text{(add)} \\
&= \frac{86,240,400}{100,000,000} && \text{(square)} \\
&= 0.862404 && \text{(divide)}
\end{aligned} \tag{41}$$

Because nearly everyone tests negative and is negative, two independent classifiers with these margins would agree on about 86% of people by chance alone.

Definition 0.15 (Cohen's kappa). For an agreement table with $p_e < 1$, **Cohen's kappa** is the agreement beyond chance, as a fraction of the largest possible agreement beyond chance, built from the overall agreement (Definition 0.8) and the chance agreement (Definition 0.14):

$$\kappa \stackrel{\text{def}}{=} \frac{p_o - p_e}{1 - p_e} \tag{42}$$

Source

2.

Example 0.18 (Cohen's kappa for the COVID-19 test). For the agreement table of Example 0.8, $p_o = 0.99$ (Example 0.9) and $p_e = 0.862404$ (Example 0.17), so Cohen's kappa (Definition 0.15) is:

$$\begin{aligned}
\kappa &= \frac{0.99 - 0.862404}{1 - 0.862404} && \text{(definition of } \kappa) \\
&= \frac{0.127596}{1 - 0.862404} && \text{(subtract in the numerator)} \\
&= \frac{0.127596}{0.137596} && \text{(subtract in the denominator)} \\
&\approx 0.927 && \text{(divide)}
\end{aligned} \tag{43}$$

```
kappa_from_table <- function(a, b, c, d) {
  n <- a + b + c + d
  p_o <- (a + d) / n
  p_e <- ((a + b) * (a + c) + (c + d) * (b + d)) / n^2
  (p_o - p_e) / (1 - p_e)
}
do.call(kappa_from_table, as.list(agree_covid))
#> [1] 0.927323
```

Example 0.19 (High agreement but low kappa). Now apply the same test to $n = 100,000$ people in a population where the prevalence is only 0.5%. Then:

- $0.005 \times 100,000 = 500$ people have COVID-19 and $100,000 - 500 = 99,500$ do not;
- with 99% sensitivity, the test is positive for $a = 0.99 \times 500 = 495$ people with COVID-19 and negative for $c = 500 - 495 = 5$;
- with 99% specificity, the false positive rate is $1 - 0.99 = 0.01$, so the test is positive for $b = 0.01 \times 99,500 = 995$ people without COVID-19 and negative for $d = 99,500 - 995 = 98,505$.

```
agree_rare <- c(a = 495, b = 995, c = 5, d = 98505)
rare_metrics <- with(as.list(agree_rare), c(
  overall = (a + d) / (a + b + c + d),
  positive = 2 * a / (2 * a + b + c),
  negative = 2 * d / (2 * d + b + c),
  kappa = kappa_from_table(a, b, c, d)
))
round(rare_metrics, 4)
#> overall positive negative kappa
#> 0.9900 0.4975 0.9949 0.4937
```

The overall agreement is still 0.99, but kappa drops to about 0.49 and the positive agreement to about 0.50, while the negative agreement stays at about 0.995. The two specific agreements show where the disagreement lies: the test and the reference agree on negatives and disagree on about half of the positive calls.

Source

High overall agreement with low kappa is the first paradox described by Feinstein and Cicchetti (1990).

Remark 0.2 (Kappa depends on prevalence). When both classifiers call one category rarely, the chance agreement p_e is close to 1, so the same overall agreement leaves little room for agreement beyond chance, and kappa can be low even when the overall agreement is high, as Example 0.19 shows. I recommend reporting the positive and negative agreements alongside kappa, so that poor agreement on the rarer category is not hidden.

Source

Feinstein and Cicchetti (1990); 1.

Feinstein, Alvan R., and Domenic V. Cicchetti. 1990. “High Agreement but Low Kappa: I. The Problems of Two Paradoxes.” *Journal of Clinical Epidemiology* 43 (6): 543–49. [https://doi.org/10.1016/0895-4356\(90\)90158-L](https://doi.org/10.1016/0895-4356(90)90158-L).