Simulation Studies

Last modified: 2026-11-04 19:04:52 (PST)

Theory often cannot give the bias, variance or coverage of an estimator when the estimator or the data distribution is complicated. A simulation study measures the same properties for any estimator a computer can run. It generates many data sets from a distribution whose parameters are known, applies the estimator to each, and summarizes the results.

This page defines the parts of a simulation study.

The data-generating process

Definition 0.1 (Data-generating process)**.** The **data-generating process (DGP)** of a simulation is the joint probability distribution F_{θ_0} from which it draws its data $X_1, \dots, X_n$, with the number of observations n and the parameter value θ_0 fixed and known:

$$X_1, \dots, X_n \sim F_{\theta_0}.$$

Example 0.1 (A normal data-generating process)**.**
Let $X_1, \dots, X_n$ be IID normal random variables with expectation $\mu = 0$ and variance $\sigma^2 = 4$, with $n = 10$. The distribution of these variables is a data-generating process, and the estimand $\sigma^2 = 4$ is known exactly, so the error of any estimator of σ^2 can be computed.

Replications and simulation studies

Definition 0.1 (Replication)**.** A **replication** of a simulation is one run of its steps: draw a data set $X_1, \dots, X_n^{(m)}$ from the data-generating process, then compute the estimate $T_m \stackrel{\text{def}}{=} \hat{\theta}\left(X_1, \dots, X_n^{(m)}\right)$, where $\hat{\theta}$ is an estimator, for $m = 1, \dots, M$. Different replications draw their data sets independently.

Example 0.1 (A replication of the normal data-generating process)**.**
For the normal data-generating process, one replication draws $n = 10$ values from the normal distribution with expectation 0 and variance 4. It then computes the estimate T_m of the variance σ^2 from those 10 values. A second replication draws 10 new values, independently of the first, and computes T_{m+1}.

Definition 0.2 (Simulation study)**.** A **simulation study** runs M replications of one data-generating process and one estimator, then summarizes the estimates $T_1, \dots, T_M$ to learn the estimator's sampling distribution.

Example 0.2 (A simulation study of a variance estimator)**.**
For the normal data-generating process, a simulation study of one estimator of the variance σ^2 runs M replications of that process and that estimator. The M estimates $T_1, \dots, T_M$ then show how the estimator behaves across repeated data sets.

References